

Q6)

Ans: DefaultAzureCredential, Create

HOTSPOT -

You have a Python application named App1 that integrates with a Microsoft Foundry project named Project1.

You need to ensure that App1 meets the following requirements:

Authenticates by using a Microsoft Entra managed identity


Sends prompts to a deployed model by using the Azure OpenAI Responses API

How should you complete the Python code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
from azure.identity import DefaultAzureCredential
from azure.ai.projects import AIProjectClient

credential =  ()

project_client = AIProjectClient(
    endpoint="https://contosoai.services.ai.azure.com/api/projects/project1",
    credential=credential,
)

with project_client.get_openai_client() as openai_client:
    response = openai_client.responses. (

        model="trail-guide-chat",
        input="Create a 3-day hiking itinerary near Seattle.",
    )

    print(response.output_text)
```

Q7).

Ans: Not(IsBlank(Local.Var01)), (Upper(Local.Var01))

You have a Microsoft Foundry project that contains a workflow for a customer support triage process.

You have an Ask a question node that stores user responses in a local variable named Var01.

You need to create the following Power Fx expressions:

An if/else condition expression that ensures that Var01 contains a value


A Send message expression that returns the stored user response in uppercase

How should you configure the expressions? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

If/else condition expression: 

Send message expression: 

Q8).

Ans:

- **Guardrails:** Select User input, Output, Tool response, and Tool call and set Action to Block.
- **Storage access:** A system-assigned managed identity that is assigned the Storage Blob Data Reader role

You have a Microsoft Foundry project that contains a customer support agent built by using the Foundry Agent Service.

The agent uploads user-provided screenshots to Azure Storage through a ticketing tool and receives a blob URL for additional reasoning.

You need to use image moderation during agent runs and prevent harmful content from being returned during runs. Azure AI Content Safety must access the images by using the blob URL. The solution must follow the principle of least privilege.

What should you configure for Content Safety? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Guardrails:

- Select Tool call and set Action to Block.
- Select User input and Output and set Action to Annotate.
- Select User input and Tool response and set Action to Annotate.
- Select User input, Output, Tool response, and Tool call and set Action to Block.

Storage access:

- Storage account access keys
- A user-assigned identity that is assigned the Storage Queue Data Contributor role
- A system-assigned managed identity that is assigned the Storage Blob Data Reader role
- A system-assigned managed identity that is assigned the Storage Blob Data Contributor role

Q9)

Ans: A

You have a Microsoft Foundry project that contains three agents as shown in the following table.

Name	Description
TriageAgent	Classifies incoming customer requests
PolicyAgent	Answers policy questions by searching internal content
ActionAgent	Creates or updates tickets by calling an HTTP API

You need to orchestrate the agents to ensure that the customer requests meet the following requirements:

Support a deterministic, step-based process that uses conditional branching and shared state across the agents.

Optionally trigger a ticket action based on the triage result.

The solution must minimize development effort.

What should you include in the solution?

- A. a workflow
- B. threads and runs without a workflow
- C. a multi-agent group chat session
- D. separate agent runs coordinated in the application code

Q10).

Ans: B

You have a Microsoft Foundry project that contains an agent. The agent uses Azure Speech in Foundry Tools.
You fine-tune a baseline speech to text model for the en-us locale and publish the model.
The agent calls the Speech to text REST API and returns an error message indicating that the project ID is invalid.
You need to set the project property to the correct ID.
To what should you set the project property?

- A. the project URL
- B. the custom speech project ID
- C. the project ID
- D. the custom speech endpoint URL

Q11).

Ans: ask_question, approval == “approved”

HOTSPOT -

You have a Microsoft Foundry project that contains an agent named PaymentAgent.
PaymentAgent includes a function tool that issues customer refunds by using an external API.
You are creating a workflow in YAML.
You need to ensure that the workflow pauses for human approval and continues with the refund step only after approval is granted.
How should you complete the workflow definition? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

```
steps:
  - id: propose_refund
    type: agent
    agent: PaymentAgent
  - id: approval
    type:
      ask_question
      basic_chat
      data_transformation
  - id: execute_refund
    type: agent
    agent: PaymentAgent
    condition:
      approval == "approved"
      propose_refund.output != null
      true
```

Q12).

Ans: C

You have an Azure Speech in Foundry Tools resource that hosts a custom speech to text model deployed to a custom endpoint. An agent uses the endpoint to perform real-time speech recognition.

You are approaching the expiration date of the custom speech to text model.

What is the expected behavior when the model expires?

- A. Speech recognition requests will return a 4xx error until a new custom model is deployed.
- B. Speech recognition requests will continue to use the expired custom model until the model is removed manually.
- C. Speech recognition requests will fall back to the most recent base model for the same locale.
- D. The custom model will be deleted automatically when the model expires.

Q13).

Ans: B

You have a Microsoft Foundry project that contains a model deployment.

You have an application that calls the deployment by using the Azure OpenAI v1 API and DefaultAzureCredential.

The developers at your company receive HTTP 403 errors when they send inference requests, even after running az login.

You need to ensure that the developers can perform model inference. The solution must follow the principle of least privilege.

Which role-based access control (RBAC) role should you assign to the developers?

- A. Cognitive Services User
- B. Cognitive Services OpenAI User
- C. Contributor
- D. Cognitive Services Data Reader

Q14). Ans: A

You have a Microsoft Foundry project that contains an agent. The agent has a Model Context Protocol (MCP) tool that queries a knowledge base stored in Azure AI Search.

Some agent runs return answers from the base model without invoking the knowledge base, which results in responses without grounded citations.

You are provided with the following code snippet that runs the agent.

```
run = project_client.agents.runs.create_and_process(  
    thread_id=thread.id,  
    agent_id=agent.id  
)
```

You need to add the correct tool _choice parameter to the code to deterministically force the agent to invoke the MCP tool on each run.

What should you add?

- A. tool_choice={"required"}
- B. tool_choice={"auto"}
- C. tool_choice={"type":"knowledge_base"}
- D. tool_choice={"type":"mcp"}

Q15).

Ans:

- **Unsupported responses:** Groundedness evaluation metrics
- **Policy violations:** Risk and safety metrics

DRAG DROP -

You have a Microsoft Foundry project that contains a customer support agent grounded in internal documentation.

After a recent update, users report the following issues:

Some answers are unsupported by retrieved documents.

A small number of responses are flagged for policy violations.

You need to evaluate each issue.

Which observability signals should you use for each issue? To answer, drag the appropriate observability signals to the correct issues. Each observability signal may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Observability signals

Groundedness evaluation metrics

Latency breakdown traces

Risk and safety metrics

Token usage analytics

Answer Area

Unsupported responses:

Observability signal

Policy violations:

Observability signal

Q16).

Ans: C

You have a Microsoft Foundry project named Project1 that contains an agent. The agent uses an OpenAPI 3.0 specification to call an external weather service.

The weather service requires a key to be passed in an HTTP header. The key value is stored as a connection in Project1.

You need to ensure that the key value from the connection is included automatically whenever the OpenAPI tool is invoked.

What should you configure in the OpenAPI specification?

- A. a header parameter defined for each operation
- B. an Azure Key Vault connection
- C. an API key security scheme
- D. a Bearer token security scheme

Q17).

Ans: B

You have a Microsoft Foundry project that serves a high-volume chat app.
Most requests are simple FAQs, but some require advanced reasoning.
You need to reduce costs and latency for common queries, without degrading the quality of the responses to complex questions.
What should you do?

- A. Route all the requests to a smaller model.
- B. Use a model cascade that routes the requests to different models.
- C. Increase the value of the max_tokens parameter for all the requests.
- D. Route all the requests to the most capable model.

Q18).

Ans:

- **Metrics to enable: Model Availability Rate and Provisioned Utilization**
- **Diagnostic log to collect: RequestResponse**

HOTSPOT -

You have a Microsoft Foundry project that contains an internal Q&A agent.

Users report the following issues when they ask the agent questions:

An increase in the following response: "No relevant information found"

Periodic HTTP 429 rate limit exceeded errors during peak hours

You need to identify whether each issue is caused by model unavailability, resource limits, or inference failures.

What should you do? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Metrics to enable:

Model Availability Rate and Provisioned Utilization
Only Tokens Cache Match Rate
Only Total Requests filtered to status code 200
Time To Response and Total Tokens

Diagnostic log to collect:

AllMetrics
audit
RequestResponse
trace

Q19)

Ans: D

You have a Microsoft Foundry project that contains a high-traffic agent.

After a recent update, operational costs increase significantly.

Monitoring confirms that the volume of user traffic to the agent remains unchanged.

You suspect that changes to the request or response characteristics are causing the increase. You need to identify whether the additional costs are driven by the model input size, the model output size, or expanded tool usage.

Which observability capability should you use?

- A. latency
- B. evaluation metrics
- C. run success rate
- D. token usage

Q20)

Ans:

- **Set tool_choice to: required**
- **Configure the tool to authenticate by: Using a distinct agent identity bound to the client application**

HOTSPOT -

You have a Microsoft Foundry project that contains an agent.

The agent uses tools to retrieve internal content and call external APIs. The agent is configured to let the model decide when to call the tools.

You need to publish the agent for a compliance workflow. The solution must meet the following requirements:

Each workflow run must include a retrieval step before generating a response.

Tool calls must authenticate by using the published agent's own identity.

Tool access must use an identity isolated from other project resources.

Tool access must use support audit tracing.

What should you do? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Set tool_choice to:

Configure the tool to authenticate by:

Q21)

Ans: D

You have a Microsoft Foundry project named Project1 that contains the following:

An OpenAPI tool that calls an external API

A project connection named Connection1 that stores the API key of the external API

When an agent calls the OpenAPI tool, the API returns a 401 unauthorized error, and traces show that the API key header is NOT being sent.

You need to ensure that the OpenAPI tool automatically includes the API key from Connection1 on all requests.

What should you do?

- A. Enable identity passthrough so that the tool uses the Microsoft Entra token of the caller.
- B. Add the API key header manually to the OpenAPI specification.
- C. Configure the tool to use the default connection of Project1.
- D. Connect the tool to Connection1.

Q22).

Ans: D

You have a Microsoft Foundry project that contains a customer support agent. The agent calls an internal knowledge API tool before generating responses.

Users report the following issues:

Some requests take more than 15 seconds to complete.

Some responses are incorrect, even when the knowledge API returns the expected data.

You need to inspect individual agent runs to view the ordered sequence of large language model (LLM) calls, tool invocations, and timing information.

Which observability capability should you use?

- A. token usage
- B. monitoring
- C. safety metrics
- D. tracing

Q23).

Ans: No

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a multimodal AI generative model that accepts image uploads and uses extracted image text to generate responses.

You discover that users can upload unsafe images and embed hidden instructions into images to manipulate the model.

You need to implement controls to mitigate the risk.

Solution: You configure a prompt shield for user prompts.

Does this meet the goal?

A. Yes

B. No

Q24).

Ans: No

You have a multimodal AI generative model that accepts image uploads and uses extracted image text to generate responses.

You discover that users can upload unsafe images and embed hidden instructions into images to manipulate the model.

You need to implement controls to mitigate the risk.

Solution: You configure image moderation to block unsafe content before processing the images.

Does this meet the goal?

A. Yes

B. No

Q25)

Ans: No

You have a multimodal AI generative model that accepts image uploads and uses extracted image text to generate responses.

You discover that users can upload unsafe images and embed hidden instructions into images to manipulate the model.

You need to implement controls to mitigate the risk.

Solution: You configure a prompt shield for documents.

Does this meet the goal?

A. Yes

B. No

Q27). Case Study:

Ans: A

Agent1 must answer questions only about the products sold by Contoso.

You need to recommend a solution to assess the responses generated by Agent1 when the agent uses the product information stored in storage1. The solution must meet the technical requirements.

What should you include in the recommendation?

- A. a Retrieval Augmented Generation (RAG) evaluator
- B. a custom guardrail
- C. model fine-tuning
- D. a groundedness evaluator

Q26).

Ans: No

You have a multimodal AI generative model that accepts image uploads and uses extracted image text to generate responses.

You discover that users can upload unsafe images and embed hidden instructions into images to manipulate the model.

You need to implement controls to mitigate the risk.

Solution: You configure protected material detection.

Does this meet the goal?

- A. Yes
- B. No

Q29).

Ans: D

You have a Microsoft Foundry project.

You plan to build a customer support solution that contains an agent. The solution must meet the following requirements:

Provide accurate, context-aware responses grounded in internal product documentation stored in Azure AI Search.

Require deep, multi-step reasoning across long contexts.

Generate detailed natural language responses.

Which type of model should you use to power the agent?

- A. a multimodal model
- B. a small language model (SLM)
- C. a key phrase extraction model
- D. a large language model (LLM)

Q30).

Ans: tool_choice: Required

DRAG DROP -

You have a Microsoft Foundry project that contains a deployed ticket-triage agent.

You discover that sometimes the agent responds without calling any tools, even when a tool is required.

You need to ensure that the agent calls a tool during execution.

How should you complete the Python code? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Values

-
-
-
-
-
-

Answer Area

```
run_payload = {  
    "assistant_id": agent_id,  
     :  ,  
    "metadata": {  
        "scenario": "ticket-triage"  
    }  
}
```

Q31).

Ans: B

You are building a web app named App1 that generates responses by using a model deployed to a Microsoft Foundry project named Project1.

Before sending the prompts to the model, App1 must retrieve documents by using Azure AI Search.

You need to integrate Project1 and App1. The solution must meet the following requirements:

Multiple client applications must use the same search configuration.

A security policy must prevent key-based authentication.

Administrative effort must be minimized.

What should you do?

- A. Create a custom HTTP connection in Foundry and manually configure Azure AI Search endpoints per application.
- B. Configure an Azure AI Search connection in Project1 and reference the connection in each application.
- C. Call Azure AI Search directly from each application by using Microsoft Entra authentication.
- D. Enable a managed identity for each application and call Azure AI Search directly.

Q32).

Ans: Grounding with bing search, code interpreter, and File search

DRAG DROP -

You have a Microsoft Foundry project that contains an agent used by the financial analysts at your company.

You need to optimize the agent workflow by providing additional data access and processing capabilities. The solution must meet the following requirements:

Ensure that the agent can perform calculations during conversations.

Ensure that the agent can access up-to-date information from public websites.

Ensure that the agent can retrieve information from documents uploaded directly to the agent.

What should you use for each requirement? To answer, drag the appropriate tools to the correct requirements. Each tool may be used once, more than once, or not at all.

You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Tools	Answer Area
☐ Code interpreter	Access up-to-date information from public websites: Tool
☐ Computer use	Perform calculations during conversations: Tool
☐ File search	Retrieve information from documents uploaded directly to the agent: Tool
☐ Grounding with Bing Search	
☐ Microsoft Fabric	

Show Suggested Answer

Q33).

Ans: B & C

You have a Microsoft Foundry project that contains a prompt agent used by a customer support web app.

The agent is invoked from a Python service that does NOT run in the Foundry portal.

You need to implement end-to-end tracing to capture latency breakdowns and exceptions across agent runs.

Which two components can you use? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. a Log Analytics workspace
- B. Application Insights
- C. OpenTelemetry
- D. the Azure Monitor Agent
- E. Microsoft Sentinel

Q34).

Ans: A

You have a customer support agent that uses the Microsoft Foundry Agent Service.

Sometimes, customers return to a session days later to continue the same support case, and the agent must resume with the full historical context. The agent must provide the following:

Multi-turn continuity within the session

Cross-session continuity for the same case

Access to the full interaction history, including user messages, agent messages, tool calls, and tool outputs

You need to ensure that the agent automatically reloads the complete history on each new turn.

What should you do?

- A. Create and reuse a conversation by storing the conversation's ID and supplying the ID on subsequent requests.
- B. Persist only the final model response stored in the client application and prepend the response to future prompts.
- C. Enable memory summarization on the agent definition to persist the context automatically.

Q35)

Ans:

- **temperature: 1**
- **output_config={"effort":: "high"}**

You have a Microsoft Foundry project that contains a deployed chat model.

You have a Python service that sends API requests to the model. The service is integrated with an automated validation system that compares generated outputs against approved response patterns.

Stakeholders report that small wording differences are causing validation mismatches.

You need to update the request parameters to improve output stability. The solution must maximize reasoning quality.

How should you complete the Python code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
message = client.messages.create(  
    model="deployment-name",  
    messages=[  
        {"role": "user", "content": "Summarize the release notes in 3 bullet points."}  
    ],  
    max_tokens=800,  
    temperature=  
    thinking={"type": "enabled"},  
    output_config={"effort":  
    }  
)
```

Q36)

Ans: A

You have a chat app in a Microsoft Foundry project and an Azure AI Search vectorized index.

You need to connect to the index to meet the following requirements:

Complex questions must retrieve information from multiple chunks.

Multi-turn conversations must influence retrieval planning.

Retrievals must run in parallel to reduce latency.

Which retrieval approach should you use?

- A. iterative retrieval
- B. agentic Retrieval Augmented Generation (RAG)
- C. chain of thought
- D. classic Retrieval Augmented Generation (RAG)

Q37).

Ans: A & C

You need to recommend a plan to create a customer support agent by using the Microsoft Foundry Agent Service. The agent must meet the following requirements:

- Retain user preferences across multiple conversations.
- Enable users to provide contextual grounding by directly uploading documents during a chat.

Which Foundry capability should you recommend for each requirement? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

To retain user preferences across conversations, use:

Agent memory that uses persistent storage
Conversation history
Orchestration-managed session context

To enable users to provide contextual grounding during chats, use the:

Azure AI Search tool
Code interpreter tool
File search tool

Q38).

Ans: A: Add a retry evaluation before the responses are returned

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

You need to improve response completeness. The solution must be implemented in the logic of the application code before responses are returned.

What should you do?

- A. Add a retry evaluation before the responses are returned.
- B. Decrease the value of the max_tokens parameter.
- C. Switch to Retrieval Augmented Generation (RAG).
- D. Replace the model with a smaller deployment.

Q39).

Ans: C

You have a customer support agent built by using the Microsoft Foundry Agent Service. The agent calls an Azure OpenAI model deployment.

During load testing, calls intermittently fail and return an HTTP 429 rate limit exceeded error.

You need to handle throttling to reduce call failures and improve reliability under load. The solution must remain within the service and model limits.

What should you do?

- A. Create a new thread and retry the calls immediately.
- B. Reduce the number of registered tools.
- C. Implement a retry policy that uses exponential backoff and jitter.
- D. Split uploaded content into smaller files.

Q40).

Ans: **Authentication method:** An Azure Login action that uses OpenID Connect (OIDC)

If the evaluation results are NOT met, configure the workflow to: Fail

You have a Microsoft Foundry project that contains an agent.

You use a GitHub Actions workflow for CI/CD.

You need to configure the workflow to automatically evaluate the agent when a pull request (PR) is created and prevent branches from merging if the evaluation results do NOT meet the defined thresholds.

How should you configure the workflow? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Authentication method:

A personal access token (PAT)

A user-assigned managed identity

An Azure Login action that uses OpenID Connect (OIDC)

If the evaluation results are **NOT** met, configure the workflow to:

Lock the target branch

Send an alert

Fail

Q41).

Ans: No

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

Users report that some responses omit required regulatory clauses, even when the clauses are present in the retrieved content.

You need to improve response completeness.

Solution: You increase the value of the max_tokens parameter.

Does this meet the goal?

A. Yes

B. No

Q42).

Ans: yes

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

Users report that some responses omit required regulatory clauses, even when the clauses are present in the retrieved content.

You need to improve response completeness.

Solution: You add a reflection pass that regenerates the response if the required clauses are missing.

Does this meet the goal?

A. Yes

B. No

Show Suggested Answer

Q43)

Ans:

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

Users report that some responses omit required regulatory clauses, even when the clauses are present in the retrieved content.

You need to improve response completeness.

Solution: You increase the value of the temperature parameter.

Does this meet the goal?

A. Yes

B. No

Q44).

Ans: No

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

Users report that some responses omit required regulatory clauses, even when the clauses are present in the retrieved content.

You need to improve response completeness.

Solution: You run an evaluation flow that scores responses for completeness and blocks responses that fall below a defined threshold.

Does this meet the goal?

- A. Yes
- B. No

Q45).

Ans: D

You are creating an image-editing workflow in a Microsoft Foundry project.

The workflow must meet the following requirements:

Ensure that background objects can be removed by applying a mask-based inpainting edit.

Preserve the original lighting and style of the edited images.

Use the built-in image editing controls, NOT a custom model.

You need to ensure that image edits apply exclusively inside the masked area.

How should you configure the workflow?

- A. Set generation mode to image_variation and provide the original image as a reference.
- B. Enable text_to_image mode and a prompt describing the desired background removal.
- C. Enable image_to_image mode and a high-strength value to regenerate the full image based on the prompt.
- D. Enable mask_inpainting and supply both the input image and a mask indicating which part of the image to modify.

Q46).

Ans: A

You have a Microsoft Foundry project that generates product marketing images from text prompts. After publishing several images, the legal team at your company identifies a competitor's logo on a sign in the background of an image. You need to remove only the logo, while preserving the rest of the image. What should you do?

- A. Apply a mask-based inpainting edit to the part of the image that contains the logo.
- B. Increase the prompt guidance strength.
- C. Modify the original prompt to exclude brand names.
- D. Rerun the prompt by using a different random seed.

Q47).

Ans: D

You have a Microsoft Foundry project that contains an agent. The agent ingests scanned PDF vendor invoices that contain tables and embedded QR codes. The agent must preserve the PDF layout in the extracted output to ensure that downstream processing can reference sections and tables. You plan to call Azure Content Understanding in Foundry Tools. You need to extract content and layout elements and detect QR codes without requiring a language model deployment. Which built-in analyzer should you use?

- A. prebuilt-documentFieldSchema
- B. prebuilt-read
- C. prebuilt-documentSearch
- D. prebuilt-layout

Q48).

Ans: A

You have a Microsoft Foundry project that contains an agent and an image generation model deployment. The agent generates original images from user-supplied product photos. You need to ensure that the generated images maintain the product identity and visual characteristics of the provided photo. What should you do?

- A. Set the input_fidelity parameter to high.
- B. Apply a groundedness detection filter.
- C. Include a prompt and input image in the request.
- D. Decrease the value of the temperature parameter.

Q49).

Ans:

- **Prompt shields action:** Set action to block.
- **Additional mitigation:** Enable Spotlighting.

You have a Microsoft Foundry project that contains an agent.
The agent accepts user-uploaded screenshots and uses a multimodal chat model.
Some screenshots contain potentially malicious embedded text.
You need to prevent a prompt injection attack and ensure that third-party content is treated as lower trust.
How should you configure prompt shields for document attacks? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

Answer Area

Prompt shields action:

- Disable the shield.
- Set action to block.
- Set action to annotate.

Additional mitigation:

- Enable Spotlighting.
- Create a custom blocklist.
- Use optical character recognition (OCR) to extract the text from the images first.

Q50).

Ans: D

You are deploying a support agent that enables users to upload photos.
You need to automatically classify uploaded images for harmful content. The solution must block content based on severity levels.
What should you do?

- A. Apply keyword scanning to optical character recognition (OCR) output by using Azure Vision in Foundry Tools.
- B. Enable prompt shields.
- C. Use blocklists.
- D. Implement image moderation.

Q51).

Ans: B

You have an app named App1 that uses a Microsoft Foundry multimodal model deployment.

App1 runs optical character recognition (OCR) on uploaded images and appends the OCR output to the prompt as additional context.

Some uploaded images contain embedded text.

You need to prevent potentially malicious instructions from being processed by the model.

What should you use?

- A. image moderation
- B. prompt shields for documents
- C. protected material text
- D. prompt shields for user prompts

Q52).

Ans: D

You have a Microsoft Foundry project that contains an agent.

You need to process mixed-format documents that contain scanned text, tables, and multicolumn layouts. The extracted content must preserve the document structure and be converted into the Markdown format for downstream reasoning.

What should you configure first?

- A. an Azure Language in Foundry Tools text analysis model deployment
- B. a generative chat completion request
- C. an Azure OpenAI Responses API call that uses a multimodal model
- D. an Azure Content Understanding in Foundry Tools analyzer

Q53).

Ans: B

You have an application that processes scanned PDF invoices. The invoices have varied layouts and include multipage tables.

You have a pipeline that uses optical character recognition (OCR) and extracts totals and invoice numbers. The results are often incorrect because the document structure is ignored.

You need to implement a solution that provides OCR, layout analysis, and template-generalizing field extraction. The solution must NOT require training a custom model.

The solution must minimize administrative effort.

What should you include in the solution?

- A. Azure Language in Foundry Tools
- B. Azure Content Understanding in Foundry Tools
- C. an Azure Machine Learning model

Q54).

Ans: A

You have a Microsoft Foundry project that contains an agent.

The agent uses a knowledge source built from documents stored in Azure Blob Storage. The documents include digitally scanned PDFs that contain multipage tables.

You have an ingestion job that extracts only plain text, causing loss of table structure, headings, and page-number metadata.

Users frequently ask questions that require the retrieval of specific table rows across the pages.

You need to configure an ingestion job for a Retrieval Augmented Generation (RAG) pipeline that performs optical character recognition (OCR) on scanned PDFs, preserves tables and headings as structure-aware chunks, and stores page-number metadata with each chunk.

How should you configure the ingestion job?

- A. Use advanced data parsing to reingest the documents.
- B. Use OCR and page-level chunking.
- C. Use page-level OCR extraction and store each page as a single chunk.
- D. Use basic parsing and fixed-size chunking.

Q55).

Ans: A

You have a Microsoft Foundry project that contains an agent.

The agent uses Azure AI Search as the retriever.

You plan to ingest PDF into an Azure AI Search index to ensure that the agent can ground responses in texts in both documents and embedded images.

Users require citations that link to the source files.

You need to ensure that during indexing, the images are extracted into a structure that can be used as input for the built-in optical character recognition (OCR) skill.

Which indexing approach should you use?

- A. an indexer to extract image data into a normalized_images collection
- B. a Shaper skill to restructure the OCR input
- C. a skillset to run the OCR skill directly against the content field of the index
- D. the outputFieldMappings parameter to write image data to a searchable field

Q57).

Ans: C

You have a Microsoft Foundry project that contains an agent.

The knowledge source for the agent is a set of scanned PDF troubleshooting guides stored in Azure Blob Storage. The guide pages contain two-column layouts and tables.

You use Azure Content Understanding in Foundry Tools to process the PDFs.

You plan to ingest the processed content into an index for Retrieval Augmented Generation (RAG) and store extracted fields for downstream automation.

Stakeholders must be able to verify where each extracted field value came from in the original PDF and route low-reliability extractions for manual review.

You need to ensure that the Content Understanding document analyzer output includes a per-field confidence score and source grounding to locations within the source document.

What should you do?

- A. Set enableSegment to true.
- B. Provide labeled samples.
- C. Enable estimateFieldSourceAndConfidence.
- D. Configure the analyzer to use generative extraction for all fields.

Q58).

Ans: D

You are building a speech processing solution in Microsoft Foundry for a customer support platform.

The platform will transcribe live phone calls, so that supervisors at your company can view call transcripts and detect issues while the calls are in progress. The call audio will arrive as a continuous stream from the telephony system.

You need to ensure that the call transcripts appear within only a few seconds of the audio stream.

What should you do?

- A. Use text to speech by using a custom neural voice.
- B. Use speech translation to generate the transcripts into multiple languages.
- C. Run a batch transcription job on recorded audio files.
- D. Use real-time speech to text to process streaming audio input.

Q59).

Ans: B

You are creating an agent workflow in a Microsoft Foundry project to support natural voice interactions.

The agent must receive continuous audio input, convert the input into text for reasoning, and then return spoken responses to a user. The workflow must meet the following requirements:

Support turn-taking dynamics, where the agent begins to generate the speech output before the user finishes speaking.

Operate with low latency to maintain conversational experience.

You need to enable both speech to text and text to speech in a real-time agent interaction.

What should you do?

- A. Use batch transcription to convert the audio input and return text responses from the agent.
- B. Use real-time speech to text for incoming audio and text to speech for agent responses.
- C. Use an embeddings model to encode the audio, and then decode the audio into text and speech.
- D. Use speech translation to convert the audio into another language and return the translated text.

Q60).

Ans: B

You have an application named App1 that uses Azure Speech in Foundry Tools to transcribe live calls.

Transcript segments often contain both English and Spanish. App1 sends each segment to Azure Translator in Foundry Tools to translate to another language.

Sometimes, mixed-language segments result in incomplete or incorrect translations.

You need to reduce translation errors. The solution must ensure that the entire transcript is translated successfully.

What should you do before sending the segments to Translator?

- A. Use document translation to translate the entire transcript as a single document.
- B. Split the mixed-language segments into single-language segments and translate each segment separately.
- C. Enable automatic language detection for the translation request.
- D. Specify English as the source language in the translation request for all the segments.

Q63).

Ans: D

You have a Microsoft Foundry project named Project1.
Project1 contains an application that processes PDF vendor invoices.
You need to configure Azure Document Intelligence in Foundry Tools to generate a Markdown output that preserves the sections and table structure of the PDFs. The solution must minimize development effort.
What should you do?

- A. Configure output=figures when you analyze the PDF.
- B. Configure content=markdown when you analyze the document.
- C. Increase the confidence threshold.
- D. Set the output_content_format=ContentFormat.MARKDOWN value.

Q64).

Ans: B

You have a Microsoft Foundry project that ingests scanned PDF invoices stored in Azure Blob Storage. Each invoice contains printed fine items and has a table-based layout.
Extracted results are stored as structured JSON and used as grounding data for an agent in a Retrieval Augmented Generation (RAG) solution.
You need to create a single analyzer that meets the following requirements:
Extracts the invoice number, invoice date, vendor name, and total amount across varying templates
Returns confidence scores so that results with confidence below 0.80 can be routed for supervisor review
What should you use?

- A. a Foundry agent that has groundedness guardrails enabled to extract invoice fields and confidence scores
- B. a custom Azure Content Understanding in Foundry Tools analyzer that defines the required fields as the extracted fields and the returned confidence scores for routing
- C. the Azure Content Understanding in Foundry Tools prebuilt-layout analyzer
- D. the Azure Content Understanding in Foundry Tools prebuilt-documentSearch analyzer and search.score from the Azure AI Search results for routing

Q65).

Ans: D

You have a Microsoft Foundry project that uses Azure AI Search to ground an agent in internal documentation. After a recent content update, users report that the agent's answers have become less accurate. You need to identify whether the retrieved content is negatively influencing the model's generated responses. Which observability signal should you review?

- A. indexer status and failure history
- B. latency breakdown traces
- C. prediction drift metrics
- D. groundedness evaluation metrics

Q68).

Ans: DefaultAzureCredential, Get

```
from azure.identity import DefaultAzureCredential
from azure.ai.projects import AIProjectClient
from azure.core.credentials import AzureKeyCredential
myEndpoint = "https://contoso.services.ai.azure.com/api/projects/project1"
project_client = AIProjectClient(
    endpoint=myEndpoint,
    credential=
    )
myAgent = "Agent1"
agent = project_client.agents. (agent_name=myAgent)

print(f "Retrieved agent: {agent.name}")
```

Q69).

Ans: A and D

You have a Microsoft Foundry project that contains a Retrieval Augmented Generation (RAG) chat solution used by customer support agents.

You are adding an automated pre-production evaluation step to a CI/CD pipeline named Pipeline1. The evaluation will run against a labeled test dataset that contains support questions and the expected grounding context.

You need to ensure that Pipeline1 fails if unsupported content or a retrieval mismatch exceeds a defined threshold:

- responses include claims not supported by the retrieved source content
- retrieved source content does not align with the labeled expected context

Which two built-in evaluators should you use in Pipeline1? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Retrieval
- B. Fluency
- C. Coherence
- D. Groundedness
- E. Response Completeness

Q70).

Ans: C

You have a Microsoft Foundry project that contains a support-ticket triage agent built by using the Foundry Agent Service.

The agent uses tool to classify the ticket type and set the ticket priority.

Sometimes, the same support case continues across multiple sessions over several days.

You need to persist state by using a durable ID to ensure that the agent can automatically reuse the full interaction history. The solution must preserve previous user messages, tool calls and tool outputs across turns and sessions.

Which runtime component should you use?

- A. output item
- B. agent
- C. conversation
- D. response

Q71).

Ans:

- Groundedness and Relevance
- Protected material

You have a Microsoft Foundry project that contains a Retrieval Augmented Generation (RAG) solution.

You need to run a pre-production evaluation by using a labeled CSV dataset that contains the query, context, response and ground truth. The evaluation must measure the following:

- Whether responses address the user query
- Whether responses are supported by the provided context
- Whether responses contain sensitive or proprietary information

Which AI quality evaluation metrics should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

To measure whether the responses are supported by the provided context and address the user query:

☐ ▼

☐ Coherence and Fluency

☐ GPT similarity and F1 score

☐ Groundedness and Relevance

☐ Groundedness and ROUGE score

To measure whether responses contain sensitive or proprietary information:

☐ ▼

☐ Hateful and unfair content

☐ Indirect attack

☐ Protected material

☐ Violent content

Q72).

Ans: A & E

You plan to configure an evaluation in Microsoft Foundry for a Retrieval Augmented Generation (RAG) chat app.

You need to provide scores for groundedness, relevance, and harmful content categories.

Which two evaluation categories can you use? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. risk and safety metrics
- B. fluency evaluator
- C. similarity evaluators
- D. AI quality (NLP) metrics
- E. AI quality (AI assisted) metrics

Q73).

Ans: A

You have a Microsoft Foundry project that contains an agent. The agent uses two tools to perform the following actions:

- Use Azure AI Search to retrieve answers from a private product documentation index.
- Use the web search tool to retrieve public information on the internet.

You need to ensure that for a specific run, the agent deterministically retrieves information only from the internet.

To what should you set tool_choice?

- A. {"type": "bing_grouding"}
- B. {"type": "azure_ai-search"}
- C. "auto"
- D. "required"

Q74).

Ans: C

You have a Microsoft Foundry project that contains a customer support agent built on a deployed chat model.

The agent responses are validated by using an automated testing system that compares generated answers to stored expected outputs. Identical prompts must return consistent response to prevent automated test failures.

You need to reduce response variability, without modifying the prompt or reducing factual accuracy.

What should you do for the model?

- A. Increase the max_tokens parameter.
- B. Remove stop sequences from the requests.
- C. Decrease the temperature parameter.
- D. Increase the temperature parameter.

Q75).

Ans: D: few shot prompting

You are developing prompts for a Microsoft Foundry project that classifies incoming support tickets by category.

You need to improve accuracy by showing the model how correct classifications look, without retaining the model or storing knowledge permanently.

Which prompt engineering approach should you use?

- A. Retrieval Augmented Generation (RAG)
- B. zero-shot learning
- C. chain of thought
- D. few-shot learning

Q76).

Ans: A: session

You have a Microsoft Foundry project that contains an agent for a customer support chat app. The agent uses a memory store and a memory search tool.

You need to ensure that the conversation history does NOT persist across separate sessions.

To what should you set the scope of the memory tool?

- A. session
- B. {{\$conversationId}}
- C. {{\$userId}}
- D. global

Q77).

Ans: Userid, Memory Tool

You have a Microsoft Foundry project that contains an agent.

You need to enable long-term memory to ensure that the agent can recall user preferences across separate conversations. Stored memories must be isolated per authenticated user without the client application manually generating user IDs.

How should you complete the Python code? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all.

NOTE: Each correct selection is worth one point.

Values

"session"
"{{\$conversationId}}"
"{{\$userId}}"
[mem_store_name]
[memory_tool]
MemorySearchTool ("support_mem_store")

Answer Area

```
from azure.ai.projects.models import MemorySearchTool,
PromptAgentDefinition
mem_store_name = "agent_mem_store"
memory_tool = MemorySearchTool
                memory_store_name=mem_store_name,
                scope=
            ),
agent_def = PromptAgentDefinition(
    model= "gpt-5.2",
    instructions= "You are a customer support assistant.",
    tools=
)
```

Q78).

Ans: C

You have a web app named App1 that processes user prompts by integrating with a Microsoft Foundry project named Project1. App1 performs the following actions:

- Sends prompts directly to a model by using the Azure OpenAI Responses API
- Invokes the Azure AI Content Safety tool by using a Foundry connection within the same request

You need to configure end-to-end visibility into each step of the request workflow.

What should you do?

- A. Enable logging by using the client SDK for Content Safety.
- B. Enable logging by using Foundry Local.
- C. Enable application tracing in Project1.
- D. Route requests through the Azure OpenAI endpoint.

Q79).

Ans:

- The sequential template that passes outputs node-by-node

- Add a Condition statement.

You have a Microsoft Foundry project that contains two agents named PolicyWriter and RskReviewer.

PolicyWriter generates draft updates for customer policies, and RiskReviewer reviews the drafts.

In the visual builder, you need to create a workflow that meets the following requirements:

- Finalizes low-risk updates without manual intervention
- Ensures predictable execution across the agents
- Requires user approval for high updates

What should you configure? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Orchestration pattern: ▼

The sequential template that passes outputs node by node

The group chat template to dynamically route control between the agents

The human-in-the-loop template that pauses execution of the workflow for input

Approval checkpoints: ▼

Add a Basic chat node.

Add a Condition statement.

Add an Ask a question node.

Q80).

Ans: D

You have a Microsoft Foundry project that contains an agent named Agent1.

Agent runs successfully, but Foundry Control Plane does NOT display values for error rates, runs, and token usage, and the Traces tab is empty.

You need to ensure that Foundry Control Plane displays the appropriate values for Agent1.

What should you do?

- A. Update Agent1 to a new version.
- B. Restart Agent from Foundry Control Plane
- C. Assign to a Log Analytics workspace to Agent1.
- D. Enable Application Insights for Agent1.

Q81).

Ans: B

You have a Microsoft Foundry project that contains an agent.

The agent uses Azure Content Understanding in Foundry Too to process vendor onboarding packets. The packs include digital PDFs that contain tables and hyperlinks.

The extracted content is indexed for search and provided to a downstream agent in the Markdown format.

You need to generate a Markdown output that has a layout and a semantic structure optimized for Retrieval Augmented Generation (RAG) workflows.

Which built-in analyzer should you use?

- A. `prebuilt-documentFieldSchema`
- B. `prebuilt-documentSearch`
- C. `prebuilt-read`
- D. `prebuilt-layout`

Q82). Ans: A

You have a Microsoft Foundry project that contains an agent.

The agent uses Azure AI Search for Retrieval Augmented Generation (RAG).

You plan to ingest and index PDF product manuals.

You need to build a solution that supports semantic similarity matching. The solution must ensure that the agent retrieves relevant data when user questions use different wording than the product manuals.

Which indexing approach should you use?

- A. vector search
- B. semantic ranking
- C. suggesters
- D. analyzers

Q83). Ans: B

You have a Microsoft Foundry agent that grounds responses from an Azure Search index that contains the following:

- Searchable text fields for product names and product codes
- A vector field that stores embeddings for product descriptions

You need to ensure that users can query the index by using the following:

- Exact product names or codes
- Natural language descriptions of the products

What should you configure?

- A. vector search only
- B. hybrid search
- C. keyword search only
- D. semantic search only

Q84).

Ans: D

You have an Azure AI Search indexer that ingests PDF policy manuals.

Client applications must display page-level citations that have bounding polygons for both text and images.

You need to add a single built-in multimodal content extraction skill to the Azure AI Search skillset. The solution must meet the following requirements:

- Provide text and image location metadata.
- Extract tables that span multiple pages.

What should you add?

- A. Document Extraction
- B. Azure Content Understanding in Foundry Tools
- C. GenAI Prompt
- D. Document Layout

Q85).

Ans: String, Generate

You have a Microsoft Foundry project that uses Azure Content Understanding in Foundry Tools to analyze marketing videos.

Video segmentation is enabled.

You need to configure an analyzer to output a generated JSON field that describes the color scheme of each video segment.

How should you configure the analyzer? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Values	Answer Area
classify	
generate	Field value type:
group	Field method:
string	
table	

Q86). Ans:

- **Field value type:** string
- **Field method:** generate

DRAG DROP

-

You have a Microsoft Foundry project that uses Azure Content Understanding in Foundry Tools to analyze marketing videos.

Video segmentation is enabled.

You need to configure an analyzer to output a generated JSON field that describes the color scheme of each video segment.

How should you configure the analyzer? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Values	Answer Area
classify	
generate	Field value type:
group	Field method:
string	
table	

Show Suggested Answer

Q87).

Ans: C

You have an invoice-processing application named App1 that uses Azure Content Understanding in Foundry Tools.

You are building a new Content Understanding pipeline named Pipeline1 that must meet the following requirements:

- Compare an invoice to its related purchase order
- Validate the invoice against static vendor contact documents
- Return a single structured output that includes discrepancy findings

You need to configure Pipeline1 and expose the pipeline as a single analyzer endpoint. What should you configure?

- A. a single-file task in standard mode that uses the vendor contract provided as an additional document during analysis.
- B. a single-file task in standard mode that uses confidence scores enabled for the extracted fields.
- C. a multiple-file task in pro mode that uses the vendor contract files as reference data
- D. a multi-file task in standard mode that uses the invoice and purchase order as input to the analyzer

Q88).

Ans: D

You have a Microsoft Foundry project that generates short promotional product videos.

After several clips are approved, reviewers notice a small watermark in the top-right corner of some videos.

You need to remove the watermark without regenerating the videos.

What should you do?

- A. Modify the original prompt to exclude watermarks.
- B. Crop the video by using the size parameter.
- C. Increase the guidance scale.
- D. Apply a mask-based inpainting edit to the affected part of the video.

Q89).

Ans: D

You have a Microsoft Foundry project that generates short promotional product videos.

After several clips are approved, reviewers notice a small watermark in the top-right corner of some videos.

You need to remove the watermark without regenerating the videos.

What should you do?

- A. Modify the original prompt to exclude watermarks.
- B. Crop the video by using the size parameter.
- C. Increase the guidance scale.
- D. Apply a mask-based inpainting edit to the affected part of the video.

Q90).

Ans: B

You are building a customer support web app named App1 in Microsoft Foundry that uses a GPT realtime model.

App1 must support:

- Live, low-latency voice conversations that use Azure OpenAI
- Streaming audio input from users and playback audio responses

You need to configure a connection method that supports real-time audio streaming in client application and targets approximately 100 ms latency.

Which connection method should you use?

- A. RTMP
- B. WebRTC
- C. SIP
- D. WebSocket

Q91).

Ans: B

You have a Microsoft Foundry project that contains an agent and uses a GitHub repository. The repository contains a YAM file named File1 that defines the evaluation settings of the agent.

You need to create a GitHub Actions workflow that runs the evaluation defined in File1 when a pull request (PR) is opened. How should you configure the workflow?

- A. Set project-endpoint to the endpoint of the project.
- B. Set evaluation-config to the path of the YAML file.
- C. Set model-deployment-name to the deployed model.
- D. Set tenant-id to the Microsoft Entra tenant ID

Q92).

Ans:

- **Managed identity scope:** Enable a system-assigned managed identity at the project level.
- **Key Vault authorization method:** Assign the Key Vault Secrets User role to the managed identity.

You have a Microsoft Foundry project that contains an agent.

The agent uses a stored access key to retrieve secrets from an Azure key vault, which violates a keyless-credentials requirement.

You need to ensure that the agent can retrieve the secrets. The solution must follow the principle of least privilege.

What should you configure? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Managed identity scope:	▼ Enable a system-assigned managed identity at the Foundry level. Enable a system-assigned managed identity at the project level. Create a service principal and store the principal's client secret.
Key Vault authorization method:	▼ Add an API key to application settings. Add a Key Vault access policy for the secrets. Assign the Key Vault Secrets User role to the managed identity.

Q93).

Ans: No, Yes, Yes

You have a Microsoft Foundry project that contains a customer support application.

You create an evaluation named Run1 that has the following configurations:

- Includes risk and safety metrics
- Includes the protected material evaluation
- Includes harmful content metrics that use a medium severity threshold

You create an evaluation named Run2 that has the following configurations:

- Includes risk and safety metrics
- Includes the protected material evaluation
- Includes harmful content metrics that use a high severity threshold

You run both evaluations against a dataset named DB1 and receive the following results:

- Content harm defect rate of Run1: 12%
- Content harm defect rate of Run2: 4%
- Protected material evaluation of Run1: 6%
- Protected material evaluation of Run2: 6%

You start a fine-tuning job by using DB1. The job fails during automatic RAI checks for multiple content harm types.

You discover that the content filtering configuration is set to high severity.

You start a fine-tuning job by using DB1. The job fails during automatic RAI checks for multiple content harm types.

You discover that the content filtering configuration is set to high severity.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
Changing the content filtering configuration to low severity will resolve the fine-tuning job issues.	☐	☐
The difference between the 12% and 4% content harm defect rate is consistent with the different severity thresholds used in Run1 and Run2.	☐	☐
The identical 6% protected material evaluation values across Run1 and Run2 indicate that this metric is unaffected by the change in the severity threshold.	☐	☐

Q94).

Ans: hierarchical Span and Tool call attributes

You have a Microsoft Foundry project that contains a multi-agent solution. The agents use tool calling to query internal systems.

You need to implement responsible AI auditing to meet the following requirements:

- Capture all the nested operations across the entire agent run.
- Record tool invocation arguments and returned results as metadata.

What should you use for each requirement? To answer, drag the appropriate options to the correct targets. Each option may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Options	Answer Area
Hierarchical spans	
A KQL query filter	Capture all the nested operations across the entire agent run:
Sampling	Record tool invocation arguments and results:
Tool call attributes	
Trace sampling policy	

Q95).

Ans:

- **HTTP 429:** Implement exponential backoff and jitter in the retry logic.
- **HTTP 400:** Split content into smaller files before uploading the files

You have a Microsoft Foundry project that contains an agent. The agent uses threads and file uploads and calls an Azure OpenAI model deployment.

During load testing, calls intermittently fail and return an HTTP 429 rate limit exceeded error. Some user uploads fail and generate an HTTP 400 file size exceeded error.

You need to mitigate the errors and reduce call failures. The solution must remain within the service and model limits.

What should you do to resolve each error? To answer, drag the appropriate actions to the correct errors. Each action may be used once, more than once or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Actions

Increase tenant-wide quotas.

Move large content to files
and use file search.

Use additional agent tools
to reduce the message size.

Implement exponential backoff
and jitter in the retry logic.

Split content into smaller files
before uploading the files.

Answer Area

HTTP 429:

HTTP 400:

Show Suggested Answer

Q96).

Ans: A

You have a Microsoft Foundry project.

You need to deploy a model from the model catalog to support a search solution for internal policy documents. The model must generate vector representations of the text in the documents and of user queries.

Which type of model should you use?

- A. an embedding model
- B. an image generation model
- C. a large language model (LLM)
- D. a small language model (SLM)

Q97).

Ans: A

You have a Microsoft Foundry project.

You need to deploy a model from the model catalog to support real-time inference. The solution must meet the following requirements:

- Use key-based authentication.
- Support real-time REST API access.
- NOT consume the vCPU quota of the virtual machines in the Azure subscription.

Which type of deployment should you use?

- A. serverless API
- B. batch
- C. self-hosted container
- D. standard

Q98).

Ans: **category:** 'AzureKeyVault' and **authType:** 'AccountManagedIdentity'

You plan to create a Microsoft Foundry project named Project1 that will contain an agent and use an Azure key vault named KV1.

You need to configure a connection from Project1 to KV1.

How should you complete the Bicep code? To answer, select the appropriate options in the answer area?

NOTE: Each correct selection is worth one point.

Answer Area

```
resource existingKeyVault 'Microsoft.KeyVault/vaults@2024-11-01'
existing = {
  name: 'KV1'
  scope: resourceGroup()
}
resource connection
'Microsoft.CognitiveServices/accounts/connections@2025-04-01-
preview' = {
  name: '${aiFoundryName}-keyvault'
  parent: aiFoundry
  properties: {
    category:
    target: existingKeyVault.id
    authType:
    isSharedToAll:
      metadata: {
        ApiType: 'Azure'
        ResourceId: existingKeyVault.id
        location: existingKeyVault.location
      }
  }
}
```

Q99).

Ans:

- **Knowledge grounding:** Configure retrieval from approved data sources.

- **Memory:** Enable agent memory that uses persistent storage.

You have a Microsoft Foundry project.

You need to create a customer support agent that meets the following requirements:

- Grounds responses only in company policy documents stored in curated repositories
- Retains customer preferences across separate chat sessions

How should you configure the agent? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Knowledge grounding: ▼

Configure revival from approved data sources.

Upload the policy documents direct to the agent.

Embed the policy documents directly into the agent instructions.

Memory: ▼

Use orchestration-managed session context.

Enable agent memory that uses persistent storage.

Retain user preferences in the state of the client application

Q100).

Ans: A

You have a Microsoft Foundry project that contains an agent. The agent generates summaries from retrieved policy documents.

You need to improve response completeness. The solution must be implemented in the logic of the application code before responses are returned.

What should you do?

- A. Add a retry evaluation before the responses are returned.
- B. Decrease the value of the temperature parameter.
- C. Increase the value of the presence_penalty parameter
- D. Replace the model with a smaller deployment.

Q101). Ans: Yes, Yes, No

You develop a test method to verify the results retrieved from a call to the Azure Vision in Foundry Tools API. The call is used to analyze the existence of company logos in images. The call returns a collection of brands named brands.

You have the following code segment:

```
for brand in image_analysis.brands:
    if brand.confidence >= 0.75:
        print(f"\nLogo of {brand.name} between
{brand.rectangle.x}, {brand.rectangle.y} and {brand.rectangle.w},
{brand.rectangle.h}")
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
The code will display the name of each detected brand with a confidence equal to or higher than 75 percent.	☐	☐
The code will display coordinates for the top-left corner of the rectangle that contains the brand logo of the displayed brands.	☐	☐
The code will display coordinates for the bottom-right corner of the rectangle that contains the brand logo of the displayed brands.	☐	☐

Q102).

Ans: C

You have an Azure subscription.

You plan to build an app that will use the Azure AI DALL-E model.

You need to deploy the model.

What should you use?

- A. the Azure SDK for Python and PowerShell cmdlets.
- B. the Azure SDK for JavaScript and Azure Machine Learning Studio.
- C. Microsoft Foundry and the Azure Command Line Interface (CLI)
- D. the Azure portal and Microsoft Graph API

Q103).

Ans: open AI and encryption

You have an Azure subscription.

You need to create a new resource that will generate fictional stores in response to user prompts. The solution must ensure that the resource uses a customer-managed key to protect data.

How should you complete the script? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
az cognitiveservices account create -n myresource -g  
myResourceGroup --kind
```

▼
AIService
LanguageAuthoring
OpenAI

▼
--api-properties
--assign-identity
--encryption

```
`{  
  "keySource": "Microsoft.KeyVault",  
  "keyVaultProperties": {  
    "keyName": "KeyName",  
    "keyVersion": "secretVersion",  
    "keyVaultUri": "https://issue23056kv.vault.azure.net/"  
  }  
}
```

Q104).

Ans: 3rd and 2nd


You have a Python application collects customer comments before posting them to a public forum.

You need to send a text comment to Azure AI Content Safety and return the self-harm severity from the response.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
def get_self_harm_severity (comment: str) -> int:
    key = os.environ["CONTENT_SAFETY_KEY"]
    endpoint = os.environ
["CONTENT_SAFETY_ENDPOINT"]
    client = ContentSafetyClient(endpoint,
AzureKeyCredential(key))
    request = 
    response = 
    result = next(
        item for item in response.categories_analysis
        if item.category == TextCategory.SELF_HARM
    )
    return result.severity
```

The first dropdown menu for 'request' contains the following options:

- AnalyzeTextOptions(categories=comment)
- AnalyzeTextOptions(text=[comment])
- AnalyzeTextOptions(text=comment)
- TextCategory.SELF_HARM(comment)

The second dropdown menu for 'response' contains the following options:

- client.analyze_image(request)
- client.analyze_text(request)
- client.moderate_text(request)
- client.path("/text:analyze").post(request)

Q105).

Ans: D

You have a custom named entity recognition (NER) project in Azure Language in Foundry Tools for support tickets. The schema for the project contains an entity type named ContactInfo.

In tagged training files, ContactInfo is used for phone numbers, email addresses, and social media handles.

Model evaluation shows low precision for ContactInfo, including false positives in which nearby text is extracted as ContactInfo.

You need to improve the precision of the project.

What should you do before retraining the model?

- A. Lower the confidence threshold for ContactInfo.
- B. Trigger an auto-labeling job.
- C. Add more support tickets as training data and label more ContactInfo entities.
- D. Replace ContactInfo by using Phone, Email, and SocialMedia entities. Relabel every matching span.

Q106).

Ans: B

You are building a text-to-speech solution that uses Azure Speech in Foundry Tools to read instructions from the script in a text file.

You discover that the solution often pronounces technical terms incorrectly.

You need to prevent the incorrect pronunciations. The solution must minimize development effort.

What should you do?

- A. From Speech Studio, train a custom neural voice
- B. Use Speech Synthesis Markup Language (SSML) to specify phonemes.
- C. Use Speech Synthesis Markup Language (SSML) to apply say as rules.
- D. Use Speech Synthesis Markup Language (SSML) to adjust the prosody of the voice.
- E. From Azure OpenAI use the Whisper model.

Q108).

Ans:

- 1. Create a project.**
- 2. Upload and tag images.**
- 3. Train the classifier model.**

You are developing an application that will detect faulty components produced on a factory production line. The components are specific to your business.

You need to use the Azure Custom Vision API to help detect common faults.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Initialize the training dataset.	
Train the classifier model.	1
Create a project.	2
Upload and tag images.	3
Train the object detection model.	

Q109).

Ans: B

You are developing a new sales system that will process user-generated video and text from a public-facing website.

You plan to notify users that their data has been processed by the sales system.

Which responsible AI principle does this help meet?

- A. fairness
- B. transparency
- C. inclusiveness
- D. reliability and safety

Q110).

Ans: C

You are designing a content management system.

You need to ensure that the reading experience is optimized for users who have reduced comprehension and learning differences, such as dyslexia. The solution must minimize development effort.

Which Azure service should you include in the solution?

- A. Azure Document Intelligence in Foundry Tools
- B. Azure Language in Foundry Tools
- C. Azure AI Immersive Reader
- D. Azure Translator in Foundry Tools

Q111).

Ans: A

You have an Azure subscription that contains an Azure App Service app named App1.

You provision a Microsoft Foundry Service resource named CSAccount1.

You need to configure App1 to access CSAccount1. The solution must minimize administrative effort.

What should you use to configure App1?

- A. the endpoint URI and subscription key
- B. the endpoint URI and an OAuth token
- C. the endpoint URI and a shared access signature (SAS) token
- D. a system assigned managed identity and an X.509 certificate

Q112).

Ans: Add a new query key, change the app to use new key, delete the compromised key

You have a web app that uses Azure AI Search.

When reviewing activity you see greater than expected search query volumes. You suspect that the query key is compromised.

You need to prevent unauthorized access to the search endpoint and ensure that users only have read only access to the documents collection. The solution must minimize app downtime.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Regenerate the primary admin key	
Regenerate the secondary admin key	1
Change the app to use the secondary admin key	2
Add a new query key	3
Change the app to use the new key	
Delete the compromised key	

Ans: Put and cognitive service

You need to create a new resource that will be used to perform sentiment analysis and optical character recognition (OCR). The solution must meet the following requirements:

- Use a single key and endpoint to access multiple services.
- Consolidate billing for future services that you might use.
- Support the use of Azure Vision in Foundry Tools in the future.

How should you complete the HTTP request to create the new resource? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area


https://management.azure.com/subscriptions/
xxxxxxxx-xxxx-xxxx-xxxx-
PATCH
POST
PUT
xxxxxxxxxxxx/resourceGroups/RG1

/providers/Microsoft.CognitiveService/accounts/CS1?api-
version=2021-04-30
{
 "location": "West US",
 "kind": " ",
 CognitiveServices
 ComputerVision
 TextAnalytics
 "sku": {
 "name": "S0"
 },
 "properties": {},
 "identity": {
 "type": "SystemAssigned"
 }
}

Q114): Ans: 0 and 25

You are building a model to detect objects in images.

The performance of the model based on training data is shown in the following exhibit.



Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Answer Area

The percentage of false positives is:

The value for the number of true positives divided by the total number of true positives and false negatives is:

Q115).

Ans: A, E, F

You are developing an application that will use Azure AI Search for internal documents.

You need to implement document-level filtering for Azure AI Search.

Which three actions should you include in the solution? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Add allowed groups to each index entry
- B. Create one index per group.
- C. Send access tokens from Microsoft Entra ID, with the search request.
- D. Retrieve all the groups.
- E. Retrieve the group memberships of the user
- F. Supply the groups as a filter for the search requests

Q116).

Ans: A

You have an Azure AI Search indexer that ingest PDF policy manuals.

Client applications must display page-level citations that have bounding polygons for both text and images.

You need to add a single built-in multimodal content extraction skill to the Azure AI Search skillset. The solution must meet the following requirements:

- Provide text and image location metadata.
- Extract tables that span multiple pages.

What should you add?

- A. Document Layout
- B. Document Extraction
- C. Azure Content Understanding
- D. GenAI Prompt

Q117).

Ans: Object projection, File projection

You are creating an enrichment pipeline that will use Azure AI Search. The knowledge store contains unstructured JSON data and the text from scanned PDF documents.

Which projection type should you use for each data type? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

JSON data: ▼

File projection

Object projection

Table projection

Extracted text data: ▼

File projection

Object projection

Table projection