

Test

Q69).

Ans: A and D

You have a Microsoft Foundry project that contains a Retrieval Augmented Generation (RAG) chat solution used by customer support agents.

You are adding an automated pre-production evaluation step to a CI/CD pipeline named Pipeline1. The evaluation will run against a labeled test dataset that contains support questions and the expected grounding context.

You need to ensure that Pipeline1 fails if unsupported content or a retrieval mismatch exceeds a defined threshold:

- responses include claims not supported by the retrieved source content
- retrieved source content does not align with the labeled expected context

Which two built-in evaluators should you use in Pipeline1? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Retrieval
- B. Fluency
- C. Coherence
- D. Groundedness
- E. Response Completeness

Q70).

Ans: C

You have a Microsoft Foundry project that contains a support-ticket triage agent built by using the Foundry Agent Service.

The agent uses tool to classify the ticket type and set the ticket priority.

Sometimes, the same support case continues across multiple sessions over several days.

You need to persist state by using a durable ID to ensure that the agent can automatically reuse the full interaction history. The solution must preserve previous user messages, tool calls and tool outputs across turns and sessions.

Which runtime component should you use?

- A. output item
- B. agent
- C. conversation
- D. response

Q71).

Ans:

- Groundedness and Relevance
- Protected material

You have a Microsoft Foundry project that contains a Retrieval Augmented Generation (RAG) solution.

You need to run a pre-production evaluation by using a labeled CSV dataset that contains the query, context, response and ground truth. The evaluation must measure the following:

- Whether responses address the user query
- Whether responses are supported by the provided context
- Whether responses contain sensitive or proprietary information

Which AI quality evaluation metrics should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

To measure whether the responses are supported by the provided context and address the user query:

Coherence and Fluency
GPT similarity and F1 score
Groundedness and Relevance
Groundedness and ROUGE score

To measure whether responses contain sensitive or proprietary information:

Hateful and unfair content
Indirect attack
Protected material
Violent content

Q72).

Ans: A & E

You plan to configure an evaluation in Microsoft Foundry for a Retrieval Augmented Generation (RAG) chat app.

You need to provide scores for groundedness, relevance, and harmful content categories.

Which two evaluation categories can you use? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. risk and safety metrics
- B. fluency evaluator
- C. similarity evaluators
- D. AI quality (NLP) metrics
- E. AI quality (AI assisted) metrics

Q73).

Ans: A

You have a Microsoft Foundry project that contains an agent. The agent uses two tools to perform the following actions:

- Use Azure AI Search to retrieve answers from a private product documentation index.
- Use the web search tool to retrieve public information on the internet.

You need to ensure that for a specific run, the agent deterministically retrieves information only from the internet.

To what should you set tool_choice?

- A. {"type": "bing_grouding"}
- B. {"type": "azure_ai-search"}
- C. "auto"
- D. "required"

Q74).

Ans: C

You have a Microsoft Foundry project that contains a customer support agent built on a deployed chat model.

The agent responses are validated by using an automated testing system that compares generated answers to stored expected outputs. Identical prompts must return consistent response to prevent automated test failures.

You need to reduce response variability, without modifying the prompt or reducing factual accuracy.

What should you do for the model?

- A. Increase the max_tokens parameter.
- B. Remove stop sequences from the requests.
- C. Decrease the temperature parameter.
- D. Increase the temperature parameter.

Q75).

Ans: D: few shot prompting

You are developing prompts for a Microsoft Foundry project that classifies incoming support tickets by category.

You need to improve accuracy by showing the model how correct classifications look, without retaining the model or storing knowledge permanently.

Which prompt engineering approach should you use?

- A. Retrieval Augmented Generation (RAG)
- B. zero-shot learning
- C. chain of thought
- D. few-shot learning

Q76).

Ans: A: session

You have a Microsoft Foundry project that contains an agent for a customer support chat app. The agent uses a memory store and a memory search tool.

You need to ensure that the conversation history does NOT persist across separate sessions.

To what should you set the scope of the memory tool?

- A. session
- B. {{\$conversationId}}
- C. {{\$userId}}
- D. global

Q77).

Ans: Userid, Memory Tool

You have a Microsoft Foundry project that contains an agent.

You need to enable long-term memory to ensure that the agent can recall user preferences across separate conversations. Stored memories must be isolated per authenticated user without the client application manually generating user IDs.

How should you complete the Python code? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all.

NOTE: Each correct selection is worth one point.

Values

"session"
"{{\$conversationId}}"
"{{\$userId}}"
[mem_store_name]
[memory_tool]
MemorySearchTool ("support_mem_store")

Answer Area

```
from azure.ai.projects.models import MemorySearchTool,
PromptAgentDefinition
mem_store_name = "agent_mem_store"
memory_tool = MemorySearchTool
                 memory_store_name=mem_store_name,
                 scope=
            ),
agent_def = PromptAgentDefinition(
    model= "gpt-5.2",
    instructions= "You are a customer support assistant.",
    tools=
)
```

Q78).

Ans: C

You have a web app named App1 that processes user prompts by integrating with a Microsoft Foundry project named Project1. App1 performs the following actions:

- Sends prompts directly to a model by using the Azure OpenAI Responses API
- Invokes the Azure AI Content Safety tool by using a Foundry connection within the same request

You need to configure end-to-end visibility into each step of the request workflow.

What should you do?

- A. Enable logging by using the client SDK for Content Safety.
- B. Enable logging by using Foundry Local.
- C. Enable application tracing in Project1.
- D. Route requests through the Azure OpenAI endpoint.

Q79).

Ans:

- The sequential template that passes outputs node-by-node

- Add a Condition statement.

You have a Microsoft Foundry project that contains two agents named PolicyWriter and RskReviewer.

PolicyWriter generates draft updates for customer policies, and RiskReviewer reviews the drafts.

In the visual builder, you need to create a workflow that meets the following requirements:

- Finalizes low-risk updates without manual intervention
- Ensures predictable execution across the agents
- Requires user approval for high updates

What should you configure? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Orchestration pattern: ▼

The sequential template that passes outputs node by node

The group chat template to dynamically route control between the agents

The human-in-the-loop template that pauses execution of the workflow for input

Approval checkpoints: ▼

Add a Basic chat node.

Add a Condition statement.

Add an Ask a question node.

Q80).

Ans: D

You have a Microsoft Foundry project that contains an agent named Agent1.

Agent runs successfully, but Foundry Control Plane does NOT display values for error rates, runs, and token usage, and the Traces tab is empty.

You need to ensure that Foundry Control Plane displays the appropriate values for Agent1.

What should you do?

- A. Update Agent1 to a new version.
- B. Restart Agent from Foundry Control Plane
- C. Assign to a Log Analytics workspace to Agent1.
- D. Enable Application Insights for Agent1.

Q81).

Ans: B